

Class 10: Structural Bioinformatics pt. 1

Ashley Allen (PID: A14633373)

Table of contents

1. The PDB database	1
2. Using Mol*	5
3. Introduction to Biod3D in R	10
4. Predicting functional dynamics	12

1. The PDB database

The main repository of bio molecular structure data is called the PDB found at <http://www.rcsb.org/>

Let's see what this database contains. I went to PDB > Analyze > PDB Statistics > By Exp method and molecular type.

```
pdbstats <- read.csv("Data Export Summary.csv")
pdbstats
```

	Molecular.Type	X.ray	EM	NMR	Multiple.methods	Neutron	Other
1	Protein (only)	169,563	16,774	12,578	208	81	32
2	Protein/Oligosaccharide	9,939	2,839	34	8	2	0
3	Protein/NA	8,801	5,062	286	7	0	0
4	Nucleic acid (only)	2,890	151	1,521	14	3	1
5	Other	170	10	33	0	0	0
6	Oligosaccharide (only)	11	0	6	1	0	4
	Total						
1	199,236						
2	12,822						
3	14,156						
4	4,580						

```
5      213
6      22
```

Q1: What percentage of structures in the PDB are solved by X-Ray and Electron Microscopy.

```
pdbstats$X.ray
```

```
[1] "169,563" "9,939"  "8,801"  "2,890"  "170"    "11"
```

The comma in these numbers is causing them to be read as characters rather than numeric. I can fix this by replacing “,” for nothing ” with the `sub()` function:

```
x <- pdbstats$X.ray
sum( as.numeric( sub(",", "", x)))
```

```
[1] 191374
```

Or I can use the **readr** package and the `read_csv()`

```
library("readr")
```

Warning: package 'readr' was built under R version 4.3.3

```
pdbstats <- read_csv("Data Export Summary.csv")
```

```
Rows: 6 Columns: 8
```

```
-- Column specification -----
```

```
Delimiter: ","
```

```
chr (1): Molecular Type
```

```
dbl (3): Multiple methods, Neutron, Other
```

```
num (4): X-ray, EM, NMR, Total
```

```
i Use `spec()` to retrieve the full column specification for this data.
```

```
i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

```
pdbstats
```

```
# A tibble: 6 x 8
  `Molecular Type`  `X-ray`    EM    NMR `Multiple methods` Neutron Other  Total
  <chr>            <dbl> <dbl> <dbl>          <dbl>   <dbl> <dbl> <dbl>
1 Protein (only)    169563 16774 12578          208     81    32 199236
2 Protein/Oligosacc~ 9939   2839   34           8       2     0  12822
3 Protein/NA        8801   5062   286          7       0     0  14156
4 Nucleic acid (onl~ 2890    151  1521         14       3     1   4580
5 Other             170     10    33           0       0     0    213
6 Oligosaccharide (~ 11      0     6            1       0     4     22
```

I want to clean the column names so they are all lowercase and don't have spaces.

```
colnames(pdbstats)
```

```
[1] "Molecular Type"  "X-ray"          "EM"             "NMR"
[5] "Multiple methods" "Neutron"        "Other"          "Total"
```

```
library("janitor")
```

Warning: package 'janitor' was built under R version 4.3.3

Attaching package: 'janitor'

The following objects are masked from 'package:stats':

```
chisq.test, fisher.test
```

```
df <- clean_names(pdbstats)
df
```

```
# A tibble: 6 x 8
  molecular_type      x_ray    em    nmr multiple_methods neutron other  total
  <chr>            <dbl> <dbl> <dbl>          <dbl>   <dbl> <dbl> <dbl>
1 Protein (only)    169563 16774 12578          208     81    32 199236
2 Protein/Oligosacchar~ 9939   2839   34           8       2     0  12822
3 Protein/NA        8801   5062   286          7       0     0  14156
4 Nucleic acid (only)  2890    151  1521         14       3     1   4580
5 Other             170     10    33           0       0     0    213
6 Oligosaccharide (onl~ 11      0     6            1       0     4     22
```

Total # of X-ray structures

```
sum(df$x_ray)
```

```
[1] 191374
```

Total # of structures

```
sum(df$total)
```

```
[1] 231029
```

% of X-ray structures

```
sum(df$x_ray) / sum(df$total) * 100
```

```
[1] 82.83549
```

% of EM structures

```
sum(df$em) / sum(df$total) * 100
```

```
[1] 10.75017
```

Q2: What proportion of structures in the PDB are protein?

```
sum(df$total[1:3]) / sum(df$total) * 100
```

```
[1] 97.91585
```

Q3: Type HIV in the PDB website search box on the home page and determine how many HIV-1 protease structures are in the current PDB? (unanswered)

When searching HIV and adding HIV-1 and protease as sub searches you get a result of 24,695 structures.

2. Using Mol*

The main Mol* homepage <https://molstar.org/viewer/> We can input our own PDB files or just give it a PDB database accession code (4 letter PDB code). We can use this markdown code to insert an image




Figure 1: Molecular view of 1HSG

Q4: Water molecules normally have 3 atoms. Why do we see just one atom per water molecule in this structure?

By having water represented only as one atom compared to three we are able to see our protein structure clearly. If water were represented how it truly is, we wouldn't be able to see our protein.

Q5: There is a critical “conserved” water molecule in the binding site. Can you identify this water molecule? What residue number does this water molecule have



> Q6: Generate and save a figure clearly showing the two distinct chains of HIV-protease along with the ligand. You might also consider showing the catalytic residues ASP 25 in each chain and the critical water (we recommend “Ball & Stick” for these side-chains). Add this figure to your Quarto document.




Figure 2: Molecular Surface of HSG



> Q7: [Optional] As you have hopefully observed HIV protease is a homodimer (i.e. it is composed of two identical chains). With the aid of the graphic display can you identify secondary structure elements that are likely to only form in the dimer rather than the monomer?




Figure 3: The important ASP25 amino-acids

3. Introduction to Biod3D in R

We can use the **bio3D** package for structural bioinformatics to read PDB data into R.

```
library(bio3d)
```

Warning: package 'bio3d' was built under R version 4.3.3

```
pdb <- read.pdb("1hsg")
```

Note: Accessing on-line PDB file

```
pdb
```

```
Call: read.pdb(file = "1hsg")
```

```
Total Models#: 1
Total Atoms#: 1686, XYZs#: 5058 Chains#: 2 (values: A B)

Protein Atoms#: 1514 (residues/Calpha atoms#: 198)
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)

Non-protein/nucleic Atoms#: 172 (residues: 128)
Non-protein/nucleic resid values: [ HOH (127), MK1 (1) ]
```

```
Protein sequence:
PQITLWQRPLVTIKIGGQLKEALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYD
QILIEICGHKAIGTVLVGPTPVNIIGRNLLTQIGCTLNFPQITLWQRPLVTIKIGGQLKE
ALLDTGADDTVLEEMSLPGRWKPKMIGGIGGFIKVRQYDQILIEICGHKAIGTVLVGPTP
VNIIGRNLLTQIGCTLNF
```

```
+ attr: atom, xyz, seqres, helix, sheet,
      calpha, remark, call
```

Q7: How many amino acid residues are there in this pdb object?

```
length(pdbseq(pdb))
```

```
[1] 198
```

Q8: Name one of the two non-protein residues?

MK1

Q9: How many protein chains are in this structure?

2 chains A and B

Looking at the `pdb` object in more details

```
attributes(pdb)
```

```
$names
[1] "atom" "xyz" "seqres" "helix" "sheet" "calpha" "remark" "call"

$class
[1] "pdb" "sse"
```

```
head(pdb$atom)
```

	type	eleno	elety	alt	resid	chain	resno	insert	x	y	z	o	b
1	ATOM	1	N	<NA>	PRO	A	1	<NA>	29.361	39.686	5.862	1	38.10
2	ATOM	2	CA	<NA>	PRO	A	1	<NA>	30.307	38.663	5.319	1	40.62
3	ATOM	3	C	<NA>	PRO	A	1	<NA>	29.760	38.071	4.022	1	42.64
4	ATOM	4	O	<NA>	PRO	A	1	<NA>	28.600	38.302	3.676	1	43.40
5	ATOM	5	CB	<NA>	PRO	A	1	<NA>	30.508	37.541	6.342	1	37.87
6	ATOM	6	CG	<NA>	PRO	A	1	<NA>	29.296	37.591	7.162	1	38.40

	segid	elesy	charge
1	<NA>	N	<NA>
2	<NA>	C	<NA>
3	<NA>	C	<NA>
4	<NA>	O	<NA>
5	<NA>	C	<NA>
6	<NA>	C	<NA>

Let's try a new function not yet in the bio3D package. It requires the **r3dmol** and **shiny** packages we need to install with `install.packages()`

```
library(r3dmol)
```

Warning: package 'r3dmol' was built under R version 4.3.3

```
library(shiny)
```

Warning: package 'shiny' was built under R version 4.3.3

```
source("https://tinyurl.com/viewpdb")  
#view.pdb(pdb, backgroundColor="black")
```

4. Predicting functional dynamics

We can use the `nma()` function in `bio3d` to predict the large scale functional motions of molecules

```
adk <- read.pdb("6s36")
```

Note: Accessing on-line PDB file
PDB has ALT records, taking A only, rm.alt=TRUE

```
adk
```

```
Call: read.pdb(file = "6s36")
```

```
Total Models#: 1
```

```
Total Atoms#: 1898, XYZs#: 5694 Chains#: 1 (values: A)
```

```
Protein Atoms#: 1654 (residues/Calpha atoms#: 214)
```

```
Nucleic acid Atoms#: 0 (residues/phosphate atoms#: 0)
```

```
Non-protein/nucleic Atoms#: 244 (residues: 244)
```

```
Non-protein/nucleic resid values: [ CL (3), HOH (238), MG (2), NA (1) ]
```

```
Protein sequence:
```

```
MRIILLGAPGAGKGTQAQFIMEKYGIPQISTGDMLRAAVKSGSELGKQAKDIMDAGKLV  
TDELVIALVKERIAQEDCRNGFLLDGFPRTIPQADAMKEAGINVDYVLEFDVPDELIVDKI  
VGRRVHAPSGRVYHVKFNPKEGKDDVTGEELTTRKDDQEETVRKRLVEYHQM TAPLIG  
YYSKEAEAGNTKYAKVDGTPVAEVRADLEKILG
```


```
+ attr: atom, xyz, seqres, helix, sheet,  
      calpha, remark, call
```

```
m <- nma(adk)
```

```
Building Hessian... Done in 0.03 seconds.
```

```
Diagonalizing Hessian... Done in 0.51 seconds.
```

```
plot(m)
```



Write out a trajectory of the predicted molecular motion:

```
mktrj(m, file="adk_m7.pdb")
```